

Lab Notes for Compressed Sensing Coding

Anthony Salerno

February 17, 2015 - July 3, 2015

Contents

0	Questions for Brian	3
0.1	Direction Based Gradient in the Objective Function	3
1	Proof of Principle (02.17.15 - 02.23.15)	4
1.1	Introduction	4
1.2	Preamble...	4
1.2.1	Hypothesis	4
1.2.2	Notes - Unpolished	4
1.3	Simulations	6
1.3.1	Preamble	6
1.4	Hypothesis	6
2	Analysis of CS Code – Sanity Check (03.10.15)	7
2.1	Preamble...	7
2.2	Down-sampling	7
3	Simulations - (03.11.15 - 04.01.15)	9
3.1	Preamble	9
3.2	Hypothesis	9
3.3	Results	9
3.3.1	Numerical Simulations	9
3.3.2	Brain Data	11
3.4	Discussion	12
3.5	Other Notes	12
3.6	Post-Meeting Notes	13
4	Directional Continuum Sampling (04.14.15)	14
4.1	Preamble	14
4.1.1	Basis	14
4.1.2	Meeting Notes	14
4.1.3	Meeting with Leigh	14
5	Return to CS with sparseMRI.v0.2 (05.01.15 - 06.03.15)	16
5.1	What happened?	16
5.2	Reanalysis and Understanding!	16
5.2.1	High Level Code	16
5.2.2	Lower Level – fnlCg.m	17
5.2.2.1	What’s Happening – Superficially	17
5.2.2.1.1	Compute the gradient of the sparse data	17
5.2.2.1.2	Compute the preobjective and objective	20
6	Editing Lustig’s Code (05.01.15 - 06.03.15)	22
6.1	Meeting With Brian (05/06/15)	22
6.2	Phase Correction	22
6.3	Weighting Function	22
6.4	Diffusion Direction based “Gradient” – Directional Redundancy	23

6.5	Problems	23
7	Continuation of Editing – Post-Vacation (06.10.15 - Current)	24
7.1	Comparison to Fully Sampled	24
7.2	What next?	25
7.2.1	Parameter Weighting	25
7.2.2	Meeting Notes – 06/24/15	25
7.2.3	Normalization – 06/29/15 - 06/30/15	26
7.3	Noise Analysis	26
7.4	Overly Smoothed Results with $\lambda_1 = \lambda_2 = 0$ – 06/30/15	27
7.5	Analysis of the code – What’s our starting point? – 07/02/15 -	28

0 Questions for Brian

Here, I will include questions that I asked Brian based on the date of the meeting – I’ll also include a basic description of the topic within the subtopics for each set of questions and whatnot.

0.1 Direction Based Gradient in the Objective Function

1. The values that we get here are somewhat dependent on how many different directions we take. It’s noted that this can be changed on the fly (which is nice) but what should be our default “threshold” of the dot product to accept?
2. The values coming from this set are on the order of 10^{14} , meanwhile, the gObj is an order of magnitude lower, and the TV and XFM are on the order of 10^9 . Do we need to normalize this value by the number of directions that we get?
3. For some reason, even though as a sum, the values should only be on the order of 10^{14} , the summation of all of these terms gives a value on the order of 10^{16} . This needs to be looked into. Regardless, this value is HIGH in comparison to our data (10^4), should this be of concern?
4. How should we change our analysis methods so that the XFM and TV weightings are comparable to the objective, and thus will actually have an effect?
5. What should we do to make the objective difference less. In the Shepp-Logan, the values are significantly lower (unless this is merely a by-product of having voxel values of $\sim 10^4$)

1 Proof of Principle (02.17.15 - 02.23.15)

1.1 Introduction

Beginning work on the “Proof of Principle” as per the 12.19.14 Meeting Notes - Plan 2

The notes state that we are to:

- Pick one slice - extract slice from dataset with FT on RO (i.e. in dim space) and no FT on PE1 and PE2 (aka PE and SL)
- Isotropic undersampling (I think this is a bad idea)
- Recon slices independently with Lustig code (using $\lambda_2 \text{psi}[m] + \lambda_1 \text{TV}[m]$)
- Add a directional similarity term to the recon (so we can recon similar slices simultaneously) ($\lambda_3 ||m_j - m_k||_2 f(\vec{d}_j \cdot \vec{d}_k)$)
- However, the form of $f((\vec{d}_j \cdot \vec{d}_k))$ is still unknown

Note that the git repository for this work can be found at <https://github.com/aasalerno/Lustig>. For pre-emptive notes on the code that Lustig wrote, please see `demo_SheppLoganTV_Notes.txt`.

1.2 Preamble...

1.2.1 Hypothesis

To begin, today we should be able to get a decent code running that can do Lustig’s code on any specified slice that we choose. I expect that the code will work as expected and provide a good rendition of the undersampled data with minimal alterations required. The hope is that we can use the basis of the code as an engine (that will need some optimization) in order to work with the data as we so hope.

1.2.2 Notes - Unpolished

First spent a ridiculous amount of time getting L^AT_EXup and running on this computer... But it works now!

The code is written both on my computer as well as on the lab computers – via a central git on git-hub – using the same datasets as required. The datasets in use are Jacob’s data, reconstructed purely using the standard recon algorithm (that is, no MATLAB involved).

For all pushes and pulls, see the git-hub repository. There are many (and it can be mapped by day!)

The code is to be written such that it is built and any values that are above 1-sampFac will be included. This means that when we are actually tacking on the r correction, we need to do it as

$$1 - \frac{r}{\max(r)}$$

Seem to have a problem with the values that I’m getting. The outer portions of k-space seem to have almost no chance of being chosen. With a penalty threshold of only 0.25 (given a sampFac of 0.5), only 40% of values are chosen. Most of the lost points are on the corners of k -space, so this may be ok... However, I need to look into **possibly making the CS type sequence specific, for things like cylindrical acquisition.**

Date: 02/20/15

As of about 11:00 (commit 6a9cfb111c508df0af9a7222ba277ef118ba147), the code `testMap.m` is up and running. It will be used in order to actually build the map of what we are undersampling and then this will be pumped into an adapted version of Lustig's code in order to get it up and running.

Encountered a bit of an issue with how the data is obtained. Since this information is only the magnitude of the data, I'm going to make a version that can handle taking the real and imaginary parts of the data, then feeding that through to the next set of functions.

As of 12:08 PM, functionality has been added for the function to handle two files containing the real and imaginary parts of the k -space information!

Now, we want to add in the CS part of the code. Here is where things become a little painful, but we can look at the datasets and actually do what is done in CS.

For some unapparent reason, the code doesn't seem to want to work. This is really annoying and irritating as the error seems to be stemming from the raw files, which may render a tonne of previous work that I've done utterly useless. Ok, scratch that – the issue was my method of plotting. I wasn't plotting the abs of the data.

As of about 5:00 it is working. I'm going to git push it and go home.

Date: 02/23/15

Upon a first look at the code that I wrote last week (adapted from `demo_SheppLoganTV.m`, the code doesn't seem to have too great of an effect. The differences between the "im_dc" (density corrected original) and the CS reconstructed data is on the order of 10^{-9} , when the data is on the order of 1.

- One of the first things that I should try is to change the TV and L1 penalties as they are currently set to 0.01 each. This may be too low to have an effect on real data – the code that I am adapting from is using a numerical phantom.
- The next step would be to try different slices or different data. Perhaps the noisiness of this data is causing a problem, but can be fixed with some more phase corrections.

I changed the values of the two penalties (λ_1 and λ_2), increasing them by a factor of 10, each becoming 0.1. The overall residual became a little bit larger, but is still on the order of 10^{-9} , which is negligible.

In order to make the data make a bit more sense (i.e. make sure I'm not making a stupid mistake on how I'm building the filter, etc.) I'm going to use Lustig's method of building the undersampling pattern.

I changed the code `testMap.m` pretty substantially, so see git-hub for a previous version. The alterations were made to try and use the other undersampling pattern. One thing that I may do later is alter Lustig's PDF that is produced by tacking on the directionality afterwards.

After changing to Lustig's method of undersampling, we see some bigger differences, but, the undersampling kind of makes it look pretty bad. It may be how the FT (using this 'p2DFT') instead of just doing an FFT as it is done in MATLAB. This seems like something weird that they don't need to do, as it just makes the data look bad... Though there must be a reason for this

As it stands – 4:00PM 02/23/15 – this is as far as I can go without further understanding how well everything is going to work for the perpendicular and parallel to the gradient direction for the undersampling technique.

1.3 Simulations

1.3.1 Preamble

Here what we want to ensure is that by applying specific filters, we are obtaining what we expect. The expectation is that using a mask that is parallel to the gradient direction (applied in k -space) is going to look significantly worse than a mask applied perpendicular to the gradient direction.

1.4 Hypothesis

Using a simulation, we will be able to tell distinct differences between the use of the two different maps. The map using a perpendicular mask on the undersampled k -space will look better than the data undersampled using the parallel mask, as the smearing will not be as bad in the cross direction of the diffusion.

The code is written in `buildLambda.m`. It is currently built just to handle a 2D case and then *eventually* handled to build a 3D. In order to make it work with the code, an idea is to make it have a z -stack (in our case, a y stack, or a ‘read-out’ stack).

2 Analysis of CS Code – Sanity Check (03.10.15)

Date: 03/10/15

2.1 Preamble...

As of this point, I've finally given my seminar (which apparently went well!) and now, I'm trying to see if there's something inherently wrong with how I'm running my CS codes, as they aren't producing the effects that I would expect. The noise looks *really* weird, and isn't how I'd expect it to be.

So, for today, we are going to change things up a bit, and move on to try and understand exactly what's happening in this code – I will try to include figures where I believe it makes sense due to the simplicity of doing so.

In order to simplify things, I made a file using the raw data from Jacob's Diffusion Weighted Imaging data taken on July 29, 2014. The raw data files are located in `/projects/muisjes/asalerno/CS/data`. Specifically we are using brain 10, direction 32. This brain was chosen to be the simplest of cases, as this is a b_0 scan, thus there is no expected major differences because of the lack of diffusion weighting.

2.2 Down-sampling

One *major* thing to note is the effect that down-sampling has. We can see it here:

Figure 1: We can see the significant blurring effect that downsampling the figure to be 128x128 has. The blurring cannot be fixed, as in order to do this, we must sacrifice some higher spatial frequency information. Note that this is using MATLAB blurring. This is explained further in the next section and figure 2.

The downsampling is absolutely required because of how the wavelet transform is done. The 2D Wavelet transform has a 2^n size requirement in order to work (Lustig 2007, Yang 2015). We don't particularly care about the number of slices being 2^n as this will be fully sampled anyway.

It seems like MATLAB does something a little weird with how it downsamples the FFT... It produces what looks like just a blurred copy of the image (as seen above) but it doesn't seem to get the data from the centre, but `instbuildLambdaead` chooses the data in the top right section (before shifting) preferentially. This can be seen in Figure 2.

Figure 2: In this figure, we see some small differences between the method of undersampling done in MATLAB vs a logical undersampling method that takes the data from the centre of k -space. All data here taken from slice 150 *after* doing a 1D FFT in the readout direction.

We can see in the above figure that MATLAB does something really weird and so, we will be using my method in order to do reconstructions that alter the number of voxels per dimension.

In order to make sure that we're using the right one, the file we must use is `kspacedS.10.32.mat`, as this is the dataset that has been undersampled properly. For this work, what we will do right now, is look at how we can try to do a reconstruction properly using compressed sensing. We may be able to do this logically using some of the existing work here. In the angiography case, they use the previous slice as the base image as “not much is expected to change”. We can do better than this, as we can try to do a combined reconstruction, where the b_0 average of the same slice (to start) would be the same.

3 Simulations - (03.11.15 - 04.01.15)

03.23.15 (work completed on and before 03.19.15)

3.1 Preamble

The point of running this simulation is to see what the best undersampling method to use would be. Our plan of attack is to do this experiment two fold, firstly with purely numerical data – i.e. a “phantom” that I design, and then afterwards using true brain data and seeing if we can get an understanding of what the best method would be.

Some of this information is noted at an earlier point in time,

3.2 Hypothesis

The expectation is that the best style would be to use the “parallel to the gradient direction” as this will give us the sharpest edge information in order to tell a cross section of the fibres that we are expecting to see. The worst should be the “perpendicular” case. It is noted that our nomenclature from before was misleading. In order to ensure that we have it correctly, the following figure explains what is meant by each.

Figure 3: We can see the direction of the gradient in the left-most panel, and the definitions of the others. The parallel will give us the “cross-fibre” information mostly, and the perpendicular will give us the “along the fibre” information.

3.3 Results

3.3.1 Numerical Simulations

For the numerical simulations, we started with a phantom that was comprised of three parts. The phantom can be seen in the following image.

When we worked with this data, we used multiple different types of undersampling filters, and compared them to the fully sampled case. The ones that we used are:

- Circle
- Square
- Parallel Strip
- Perpendicular Strip
- Parallel Ellipse

Figure 4: This figure is the numerical phantom that we used. Region (A) represents an area of all zero data, used to serve as a reference point when plotting the figure (it forces the dynamic range to begin at zero). Region (B) represents an area of isotropic diffusion, where $\lambda_1 = \lambda_2 = \lambda_3 = 10^{-3}$. This area should have an FA value of zero. Region (C) represents an area with restricted diffusion, where $\lambda_1 = 10^{-3}$, but $\lambda_2 = \lambda_3 = 5 \times 10^{-4}$. This gives a theoretical FA value of approx 0.41. It is noted that the b value used in this experiment was $1917 \frac{\text{s}}{\text{mm}^2}$

- Perpendicular Ellipse

For the two ellipses, it is noted that the long axis to short axis ratio was set to be 2 (i.e. $a/b = 2$ if $(\frac{x}{a})^2 + (\frac{y}{b})^2 = 1$). It is also noted that x and y here are those of the axis rotated relative the gradient vector direction.

Upon comparing all of our datasets, we found that the square and the circle gave the smallest RMS error. This can be seen via a bar chart – this chart combined all the ROIs from all directions (doing the fully sampled - undersampled), and and ran an RMS on all of the data.

For these plots, we used a ROI that just covered the circle – quite tightly – so that we didn't have much of a confounding factor. We still saw, though, that the “circle in the centre” (as seen in the next plot) was the best method to reduce the error that we get in our data.

Figure 5: This figure shows us how the RMS differs for the different undersampling types. We can see the undersampling types in the next figure.

Figure 6: The different filters that we used for the plots. It is noted that the codes used almost no roll off. For this example, we are assuming the gradient direction to be totally vertical.

3.3.2 Brain Data

When proceeding to brain data, I expect that we would be the same results. Using data that can be found in `/projects/muisjes/asalerno/CS/data`, and undersampling methods that can be found in `/micehome/asalerno/Documents/CompressedSensing`, we ran the same undersampling techniques, but got some interesting results. The ellipse method is now *worse* than the strip method! Reasons for this are currently unknown. The dataset being shown is brain 2, all 30

directions, however, this can (and will) be redone for more brains in order to get a more realistic understanding of what’s happening.

Figure 7: RMS error from fully sampled data in comparison to the fully sampled case.

3.4 Discussion

03.25.15 - 04.01.15 Surprisingly, we didn’t particularly get what we expected. When running the simulation for the numerical phantom, the data didn’t seem very different at all – at least when comparing FA values. The data actually gave us results that were confounding... When being compared to the “standard” methods of undersampling (i.e. Circular or low resolution undersampling), the parallel and perpendicular cases fell through and were not effective. However, even when looking for a “happy medium” – an elliptical technique, which we expect would have the lower RMS error of the circular technique, but the specificity of the directionally specific techniques – we don’t find that (with $n = 4$ for brains tested).

3.5 Other Notes

Due to space constraints, I’ve had to delete a lot of the (extraneous) data from the folders. I removed all of Jacob’s recons (as they’re unimportant, and stored elsewhere). I also deleted (inadvertently) the FA map that I produced for it, but all in all, it wasn’t that important anyway. Since it is so easy to reproduce from the original information (i.e. the fully sampled case) I’m also getting rid of the undersampled real and imaginary data.

3.6 Post-Meeting Notes

Meeting Date: 04.01.15

Brian stated that I should do some work on looking at the line profiles of everything that has been done because RMS is akin to looking primarily at DC information. This should give us a better rendition of what we want because we're interested in the edges more than the DC information.

4 Directional Continuum Sampling (04.14.15)

04.14.15 – Never completed

4.1 Preamble

4.1.1 Basis

The idea behind what this method of acquisition is that we won't necessarily need CS for a reconstruction. The notion is to use a continuum of directions, as opposed to a specific number of discrete directions – while only taking a small number of points for each “direction” we would be able to do a full reconstruction using a solid angle of points (and possibly, be able to shrink the required number of points by utilizing CS).

4.1.2 Meeting Notes

The notes are about building the simulation, as well as some work in relation to how to build the simulation. As of 04.14.15, I don't have a firm grasp of what I want to do and/or how I plan to do it. However, I believe that looking at PE tables (and including that in my solutions) will be a good method.

- If CS is to be used, incorporate the directional information before we look at 3D recon. A possible form could be:

$$e^{-\frac{\alpha_{jk}^2}{2\sigma^2}}$$

Where $\alpha_{jk} = \angle \vec{d}_j \vec{d}_k$

- Do a weighted average to obtain the proper eigenvalues for the simulation

4.1.3 Meeting with Leigh

After meeting with Leigh SN to discuss methods of how to best do something that would make sense with respect to the FSE acquisition method. She pointed me to an email from Brian and the locations of some tables.

Email from Brian

Hello Jacob et al,

You can try the following command for generation of cylindrical FSE tables.

```
/home/bjnieman/source/vnmr/cyltable_gen.py--nf64--cylindrical--etl8--etlorder2,1,0,3,4,5,6,7--angorder256256test_nf64_etl8_256_256
```

Obviously, you can replace ‘‘nf 64”, ‘‘etl 8”, ‘‘etlorder 2,1,0,3,4,5,6,7” and ‘‘256 256” (as well as the output filename) with numbers of your choice. nf should be divisible by etl.

To view the tables and inspect for problems, you can use:

```
/home/bjnieman/source/vnmr/petable_display.pytest_nf64_etl8_256_256_nf64_ni64
```

Options `-nacq` and `-etl <number>` allow you to display the number of acquisitions (repeated in the table) or the echo number of each acquisition.

Come to me if you have questions.

Brian

Email from Leigh

Here is the location of several different petables at 75 um (272 x 272 in the phase encode, *xz* plane):

```
/home/leigh/cyltables/75um/angorder/LSN_E3_angorder_272_etl8_nf64_ni80
/home/leigh/cyltables/75um/cylaxisdist/LSN_E3_cylaxisdist_272_etl8_nf64_ni80
/home/leigh/cyltables/75um/pseudocyl/LSN_E3_pseudocyl_etl8_272_nf64_ni80
```

If you need to look at all of the smaller tables (12 each) have a look in the parent folder, but know that there are a LOT of other tables in those folders, so just be careful to read the names carefully (they are different in the E3 E4 parts, where E3 means the third echo in the centre and E4 means the fourth).

To make these tables I used the following commands:

```
/home/bjnieman/source/vnmr/cyltable_gen.py--nf64--etl8--cylindrical--angorder--individ_
table_limit5120--etlorder2,1,0,3,4,5,6,7,8272272--clobberLSN_E3_angorder_272_etl8
/home/bjnieman/source/vnmr/cyltable_gen.py--nf64--etl8--cylindrical--cylaxisdist--individ_
table_limit5120--etlorder2,1,0,3,4,5,6,7,8272272--cylaxisangle90--clobberLSN_E3_cylaxisdist_
272_etl8
```

NOTE, PSEUDOCYL IS A LITTLE DIFFERENT

```
/home/bjnieman/source/vnmr/cyltable_gen.py--nf64--etl8--cylindrical--pseudocyl_pe2--individ_
table_limit5120--k0echo3--clobber272272LSN_E3_pseudocyl_etl8_272
```

Also, note that `nf` and the table axis (272) must be divisible by `etl`. These tables all have a resolution of 75 um.

Let me know if you need anything else.

Leigh

These locations as listed above will be used in order to find the information about the PE tables so I make sure that what I do is compatible with the scanner – thus I’m not testing something that can’t be done.

5 Return to CS with sparseMRI_v0.2 (05.01.15 - 06.03.15)

5.1 What happened?

When I was working with the code, at some point in time, I did something that fatally affected the code, causing problems and iterative work that would lead to results that were terrible – equivalent to doing nothing at all. The CS reconstruction looked just like the zero-filled reconstruction. To figure this out, I re downloaded the codes and reran them, as I was having problems with the phantom reconstruction as well (which should work absolutely perfectly!). The code was re downloaded and placed in `/micehome/asalerno/Documents/Lustig`.

5.2 Reanalysis and Understanding!

Here, I will discuss each code and exactly what's happening. The basis of the code will be from the `demo_SheppLoganTV.m`, just like the previous time that I did this (which I will be looking at as a reference as well!), though this time, I will *not* edit the code as much so that I don't break it again!

5.2.1 High Level Code

Note that the folder `Orthogonal` needed to be added from Lustig's other code – specifically the code that goes through Wavelets and their uses etc.

p2DFT:

I don't quite get what the point of this is, though it seems like it's some sort of 2D Discrete Fourier Transform, in matrix form. This plays off of Brian's idea of using a matrix of matrices on the diagonal to apply this type of transform to the data.

XFM:

This is the all important sparsifying transform. As a default it is set to be 1 for the phantom. For testing purposes, I will change this to be a wavelet. When this is run, using a Daubechies wavelet, at 6th order. As defined by Lustig, this means:

```
The Daubechies filters are minimal phase filters that generate wavelets which
have a minimal support for a given number of vanishing moments. They are
indexed by their length, Par, which may be one of 4, 6, 8, 10, 12, 14, 16,
18 or 20. The number of vanishing moments is par/2.
```

Using the Wavelet transform, from a standard Shepp-Logan phantom, we get decent results!

Eight Iteration for loop

Weirdly enough, this looks better as the iterations go on... It isn't running the *exact* same thing each time, but using the previous data (seemingly) for each iteration. So, in theory, having a max of 64 iterations would be the same as having 8 iterations in the `for` loop with each loop having a max iteration limit of 8. This will be tested at some point – go to it here.

Questions

1. What exactly does the p2DFT do? Why use this instead of an FFT?
2. What effects would different XFMs have (even using the same type, but with different numbers)?
3. How do the two parameters that we have (soon to be made 3) affect the image?

5.2.2 Lower Level – fnlCG.m

The prime importance of this part of the code is to understand `fnlCG.m` and the objective functions, because this is the basis of how everything is done in this code. By understanding this, I'll be able to build on it (or possibly even port it!).

Originally, this code had all of the codes in one script, however, I moved all extra functions into their own scripts, located in the folder `./grads`.

For this analysis, I will be switching between mathmode and verbatim. To make it easier, here are the switches:

Represents	Code	Math Mode
Raw Data (k -space)	<code>params.data</code>	y
Iterative Data (sparse)	<code>x</code>	x or m
Fourier Transform	<code>params.FT</code>	$\mathcal{F}_u[\cdot]$
Sparsifying Transform	<code>params.XFM</code>	$\phi[\cdot]$
Inverse Fourier Transform	<code>params.FT'</code>	$\mathcal{F}_u^{-1}[\cdot]$
Inverse Sparsifying Transform	<code>params.XFM'</code>	$\phi^{-1}[\cdot]$
Gradient of the data	<code>dx</code>	∇m

Table 1: Table of conversions between math mode and code.

5.2.2.1 What's Happening – Superficially

The first thing that is done is looking at the line search parameters. These are parameters chosen for the minimization technique that's employed here. Nothing of major importance here since it is all set previously in the `params` variable.

5.2.2.1.1 Compute the gradient of the sparse data

The equation that we're editing is: $g_0 = \nabla \phi(x)$ This calculates the gradient of the *sparse* dataset. It uses `wGradient.m`, which relies on calculation of the objective functions (which still needs to be discussed).

gOBJ.m - Gradient of the Objective Function

This code is just one line:

```
gradObj = 2*(params.XFM*(params.FT'*(params.FT*(params.XFM'*x) - params.data)));.
```

It is important to note that when the transform matrices are multiplied, it applies the transform, but multiplying by the transpose applies the inverse transform.

The important players here are:

- **XFM** - The sparsifying transform
- **FT** - The Fourier Transform
- **x** - Our data that is undergoing iterations
- **data** - The original set of data

This line does the following:

1. Returns the data to image space
2. Puts the data in k -space
3. Subtracts the original dataset from it in k -space
4. Returns the data to image space
5. Puts image into sparse space
6. Multiplies it by 2

There's some intrinsic penalty right off the bat because of how the calculations are done. Firstly, with sharp edges, the FT's are not the exact same. In addition, the sparsifying transform likely has an effect.

When looking at the phantom data, it seems that the major differences between **x** and **params.data** (not surprisingly) come from the edges, but the differences are decently sized (~ 0.5 , compared to an image max of ~ 1 .)

This code seems to do this part of the Compressed Sensing equations:

$$\|\mathcal{F}_u m - y\|_2 < \epsilon$$

Since we are taking the **abs** of this when plotting (and looking at it's value in comparison to ϵ) this is effectively taking the ℓ_2 norm of the data.

gXFM.m - Gradient of the ℓ_1 Transform Operator

This function is also a one liner:

```
grad = p*x.*(x.*conj(x)+params.l1Smooth).^ (p/2-1);
```

Here, **p** represents the pNorm value for the ℓ_1 term. It is a *scalar* value!

In mathematical notation (to be simpler to read) we have:

$$\nabla f = px(|x|^2 + \alpha_1)^{\frac{p}{2}-1}$$

In this case, again, x represents the sparsified data (i.e. it is in a sparse space) that continuously undergoes iterations. It's noted that $|x|^2 = xx^*$. α_1 here represents an ℓ_1 smoothing term. In the example here, it is set to be 10^{-15} , and p is set to be 1.

This equation doesn't seem to come up in the paper. If we analytically look at what this gives us, we can easily see what effect we should get. We will use the parameters as stated in the previous paragraph.

$$\begin{aligned}
 \nabla f &= px(|x|^2 + \alpha_1)^{\frac{p}{2}-1} & |x|^2 + \alpha_1 &\approx |x|^2 \\
 &= x(|x|^2)^{\frac{1}{2}-1} \\
 &= x(|x|^2)^{-\frac{1}{2}} \\
 &= x(|x|^{-1}) \\
 &= 1
 \end{aligned}$$

When we run the code, the difference between the range of values are 1 ± 10^{-9} . So we can see that the analysis makes sense. I will note, though, that I *don't* get what this part does exactly. Brian and I *NEED* to talk about this!

gTV.m - Gradient of the TV operator

As the title says, this handles the TV operator. This, however, is more than one line, as the TV is applied to the data in *image space*, but x is iterated in the *sparse space*.

Something of interest is that when the TV operator is applied, the size of the matrix doubles, and increases a dimension! x went from being $[256 \times 256]$ to $[256 \times 256 \times 2]$. There seems to be some orthogonality to it, as seen in Figure 8.

Figure 8: Comparison of the two $[256 \times 256]$ planes given in after the TV operator is applied

Overall, however, this equation has a similar form to the one before it, but requires a few transforms in order to be put back into the sparse space.

Let $D_x = TV[\phi^{-1}[x]]$:

$$\nabla f = \phi[TV^{-1}[pD_x(|D_x|^2 + \alpha_1)^{\frac{p}{2}-1}]]$$

The terms inside of $\phi[TV^{-1}[\dots]]$ are analogous to the terms in the above section. Again, we utilized $|D_x|^2 = D_x D_x^*$. In this case, however, it is noted that the data doesn't follow the same "1" value as the gXFM did, even though the terms inside the transforms are all 1. My guess is that the noise of this distribution (which will happen in normal DAQ) will cause this value to increase

dramatically (which is logical because variation will occur due to noise from voxel to voxel – if it didn't, the assumption of Gaussian or Rician noise would be incorrect).

5.2.2.1.2 Compute the preobjective and objective

The preobjective (and objective) functions are calculated early in order to “make it cheap to compute many times for efficient line-search”. Some of these values are calculated many times and often, so this would be a good place to attack when we are looking at optimization.

preobjective.m - Precalculates Transforms to Make the Line Search Cheap

This code is quite simple, as it just applies the transforms to both x and dx and places them in variables. The good news is that it makes them readily accessible, but the bad news is that in large datasets, this can become quite memory consuming.

The nomenclature is long, annoying, tedious, but all in all, understandable. The key is to look for each set of transforms. From this code, the outputs that we get are:

- FTXFMtx - k -space of x ($\mathcal{F}_u[\phi^{-1}[x]]$)
- FTXFMtdx - k -space of dx ($\mathcal{F}_u[\phi^{-1}[dx]]$)
- DXFMtx - TV of the image space of x ($TV[\phi^{-1}[x]]$)
- DXFMtdx - TV of the image space of dx ($TV\phi^{-1}[dx]$)

objective.m - Calculates the Objective Function

The first thing that caught me in this code is the parameter t . It's just set to 1 in this code, but this adds *yet another parameter* for us to play with.

In this code, we calculate the objective as:

```
obj = FTXFMtx + t*FTXFMtdx - params.data;
```

In mathematical terms, we have:

$$Obj = \sum_{i=1}^N (\mathcal{F}_u[\phi^{-1}[x]]_i + t\mathcal{F}_u[\phi^{-1}[dx]]_i - y_i)^2$$

Where y is the original data that we started with (off the scanner). We let i represent the i^{th} voxel of the data, since the objective value that we obtain is to be a measure of our loss. To put it into perspective, we are figuring out what the ℓ_2^2 in k -space between the data, plus it's gradient, from the original data off the scanner. For all conversions, remember to see Table 1.

This then happens again for the XFM and TV differences, but in a slightly different manner, which will be explained below.

For those two, the code is as follows:

```
w = x(:) + t*dx(:);
XFM = (w.*conj(w)+params.l1Smooth).^(p/2);
```

The only difference that exists between the two is what the w is calculated from. In TV it uses DXFMt(d)x, and in XFM it uses (d)x. This makes sense because the total variation uses the TV of image space (which is what DXFMtx is), and the XFM is calculated via the sparsified data, contained in x .

The code was written in such a way that the TV and XFM weights don't need to be constant. The overall value given from the objective function is calculated as:

```
res = obj + (TV) + (XFM) ;
```

Where the TV and XFM were multiplied by their respective weights. The RMS value is also calculated here, but it isn't used in `fnlCg.m`.

This objective function is calculated once with $t = 0$, as a reference, and then again with t as a parameter – the reference is used to break a loop. With $t = 0$, we lose all effects of the gradient, and so this is the differences that we would have with just the transforms (i.e. the problems I was talking about earlier because of the sharp edges and cutoffs, etc.).

The importance of α and β

```
f1 > f0 - alpha*t*abs(g0(:)'+dx(:))
t = t*beta
```

$f1$ is the value from the objective function when $t \neq 0$ $f0$ is the value from the objective function when $t = 0$.

α and β are very important in this analysis, as they give different weights when attempting to get out of the line search.

α tells us how strongly we will weight the effects of the gradient to take off of the original dataset if we wish to get comparable objective functions between the original dataset and the iterated dataset.

β tells us how much are we going to weight the gradient when doing our calculations for the objective function. As we weight it less and less, we *should* be getting closer and closer to the case where $t = 0$, giving us smaller values.

Once we get through the while loop via α and β we update our data (x) and our gradients (g_0) and go through the calculations again with new starting criteria unless the ℓ_2 of the gradient is less than ϵ , or we've reached our iteration limit.

6 Editing Lustig's Code (05.01.15 - 06.03.15)

For all of this work, the code has been moved to `CompressedSensing` – anything from the aforementioned `CompressedSensing` folder has been moved to `MincMatlab` for a better naming scheme. A “safe version” of the code can be found in `Lustig/sparseMRI_v0.2`.

For all of this code, the expectation is that from the data, we will have it separated by slice in the readout direction and then stacked by direction. So in a (preferably `.mat`) file, we would have one brain, one readout slice, all directions. Note that (for obvious reasons), the order of the directions should be the same as in the Gradient Vector file.

Currently using `for` loops, which is relatively slow – but will work to check brute force calculations for right now. I'll be doing this a la John Camarack, and optimizing later – `mtimesx` may be a good place to look for optimization routines...

6.1 Meeting With Brian (05/06/15)

Brian has some ideas for what needs to be changed. It can be looked at as a few points.

- Weighting Function
- Phase Correction
 - Undersampled recon and smooth it out and look at the phase $\theta = \tan^{-1} \left(\frac{\text{Im}(x_{ij})}{\text{Re}(x_{ij})} \right)$
 - Generally, the phase correction that Leigh will apply should be done before this (this one makes sure that the phase is generally okay before we do anything else)
- Make this work for multiple directions – i.e. adapt the code to work in 3D

6.2 Phase Correction

As per Brian, the way to calculate the phase would be to take the image (as it is real and imaginary) and do:

$$P = \frac{s^*}{|s|}$$

Where s represents the signal of the image. This is akin to multiplying by $e^{-\phi(s)}$.

When I run the calculation using the above equation, it looks *nothing* like what Lustig gets through his method. So, I want to talk this through with Brian, and see exactly what the best method here – is what I'm doing correct?! Calculations for the phase can either be done Brian's or Lustig's way – and can be found in `phCalc.m`. Currently, the code will do the phase calculation either way, based on `rl`, a parameter that is fed into it.

I ran a quick check, testing the phase for different slices, and determined that (with the exception of the areas of zero-padding) the phases are different.

Via a quick reality check, I made a file `brain.6.01.mat` that is zero-padded, but only the first direction (as opposed to all directions in `brain6.mat`).

6.3 Weighting Function

We want the code to be able to work on all 30 directions simultaneously and have them be weighted according to how similar the gradient directions are. In order to make sure that this is absolutely

viable, some minor edits had to be made to the `GradientVector.txt` file in order to make the magnitude of each vector as close to 1 as we can get them (currently, they are correct to 1 ± 0.004).

The weighting function that we are planning on using has the form:

$$\alpha = e^{-\frac{(\vec{d}_i \cdot \vec{d}_j)^2 - 1}{2\sigma^2}}$$

A code for this is written, it is called `dotThresh.m`. It can be found in `./diffusion`.

6.4 Diffusion Direction based “Gradient” – Directional Redundancy

Here is where we look at the actual directional redundancy that we get when using the many different (and similar) directions.

All of the gradient codes were edited to work with all 30 (or however amount) of directions that we have. Again, the expectation is that the brains are split into a different file for each slice about the readout, and each stacked image is a different direction.

6.5 Problems

Noticed that when running the objective function, my error is astronomical ($\sim 10^{13}$), so I decided to do a check with Lustig’s and his is MUCH lower. So I’m doing some reality checking with the first direction. Essentially, I ran the Lustig’s *vanilla* code with this slice, and am comparing it to the results of my altered version.

When running the two, the FT’s are identical, as are the XFM’s (as they are well defined). The error, I believe, must come from the TV...

Nope. In fact, the major difference is occurring with the directional weighting (duh!) because it is actually comparable in size to `gradObj`, ($\sim 10^2 - 10^3$), while the other two are much smaller ($\sim 10^{-2} - 10^{-1}$).

Turns out I made a stupid mistake and was dividing by the PDF in the `res` calculation! Edited this, and now all is well! Need to be more careful.

What I’ve noticed is that the directional differences are significantly more varied than the others. We have differences varying from 10^{-33} to 10^{22} . Could this be due to very different images, or possibly due to relatively different dot products (i.e. $\vec{d}_i \cdot \vec{d}_j \approx 0.25$), but the weighting isn’t dropping fast enough.

The code is SLOW. When I run a 2D test in comparison to the original, it takes about 2 seconds longer. The 3D version takes very long. I’m testing it in comparison to looping the 2D case.

7 Continuation of Editing – Post-Vacation (06.10.15 - Current)

Took a vacation and have been slacking with regards to entering my notes in here. As a quick update as to what’s going on, I have the following:

The code is up and running – right now we are doing some parameter optimization with respect to $\lambda_{1,2,3}$. This took much longer than anticipated because of some errors that I made with respect to the data. The data wasn’t FFT’d properly, causing there to be a “chequerboard” pattern on it – that is, all high frequency information. This was causing the sparsity constraint to fail miserably, as all of the information after the wavelet was in the bottom right corner, instead of the top left (this is because the wavelets are actually only working on the real part of the data).

As of now, everything is looking a bit better, however, there still does seem to be this soft roll in the data... Could this possibly be due to a misplacing of “zero” in the transition from **vrecon** to MATLAB?

7.1 Comparison to Fully Sampled

For this section, the comparison that was done was between the data that comes out of the algorithm, and the original fully sampled data. It is noted that both sets of data have individual voxel values on the order of 10^4 . The overall RMS difference that we get between the datasets (for a 256×256 dataset) is also on the order of $10^3 - 10^4$.

What I’ve found (which as of right now needs to be confirmed), is that there are sections that are slightly better than others – though not by much. These sections come up in sort of “waves” within the parameter optimization map. Below is an optimization map, done using a logarithmic map for **xfmWeights** and **TVWeights**. We can see that it seems like there are jump discontinuities, but these are not as big as they seem in Figure 9. The dynamic range of the colour axis only spans a value of less than 200, for an RMS average of about 8×10^3 , or about 10^{-1} for each voxel! In Figure 10, we can see a similar plot, but the XFM and TV values are much higher. This leads to a slight increase in the RMS, but the XFM and TV differences in the objective function are comparable to the OBJ (the ℓ_2 norm) of the data consistency term. This is what we’re looking to get, as it’s what is seen in the Shepp-Logan example.

Figure 9: Parameter optimization using an RMS difference. These values are computed using TV and XFM $\in [10^{-3}, 1]$. Classified as the “low” parameter values.

Figure 10: Parameter optimization using an RMS difference. These values are computed using TV and XFM $\in [1, 10^4]$. Classified as the “high” parameter values.

7.2 What next?

After looking at all of the figures, what I’ve found is that they don’t seem to be very different from each other if we look at the magnitude of the images. It is noted that they do look pretty nice, and in my eyes better than the merely undersampled case. I want to analyse a few things after looking at these images:

- Can we even more grossly undersample in order to see weirder effects?
- How different are the images actually? Is one of the parameters better to change for differences? Were the values that I chose to work with too close together (i.e. it was taking $\lambda_{1,2} \approx 0$)?
- What exactly is the distribution of the noise? We’ve added in incoherence noise, and removed some noise via denoising steps in the algorithm.

7.2.1 Parameter Weighting

Something that I noticed when running the same analysis on the Shepp-Logan phantom was that the XFM and TV weightings are actually close to the objective weighting, \pm an order of magnitude. This is integral information to have, because of how it compares to the case of our data, where it is many orders of magnitude lower. Perhaps we are not using the right orders of values for our analysis – this may be why we see such minuscule differences when we change our parameters.

The data for each section can be found easily. The low parameter data can be found at `/projects/muisjes/asalerno/CS/data/LowTV-XFM.mat`, and the high data can be found at `/projects/muisjes/asalerno/CS/data/HighTV-XFM.mat`. Each of these is a 4D file, where the 3rd dimension is TV and the 4th dimension is XFM weights. In each dataset, the `xfmWeights` and `TVWeights` variables can be found. This is useful if you are using something like `logim.m` to plot.

7.2.2 Meeting Notes – 06/24/15

During my meeting, Brian suggested that I change how the images are input into the algorithm. The benefit of this is that “no matter what we feed into the code, the parameters should still be the same”.

Brian also wants me to normalize the data so that either the images' maxima are set to be 1, or that $(k_x, k_y) = (0, 0)$ has a value of 1. This also means that $\bar{im} = 1$.

Eventually, Brian wants to fix this value so that we don't have this drop in intensity when applying CS to diffusion data. An idea is to either:

- Add in another term of the form $(F_o - \int f(x) dx)^2$ where $f(x)$ is the image, and F_o is the original average.
- Possibly tack on a weighting term: $\|w(\mathcal{F}_u\{f(x)\} - y)\|_2$

Later on, an idea may be to work in k -space! Problem with this, however, is that we would have to properly apply $\mathcal{F}\nabla f(x)$. Again, here, $f(x)$ is the image.

Lastly, with respect to the noise analysis (talked about later, in Section 7.3) I was able to show that the noise looked Rayleigh (and thus likely Rician). However, Brian said to look into this more by analyzing the \mathbb{R} and \mathbb{I} parts individually to check if they're both Gaussian.

7.2.3 Normalization – 06/29/15 - 06/30/15

I did a quick test of this, and it seems like by doing this, we can get the more realistic (and noted in literature) values of $\lambda_{1,2} \leq 1$. However, we still hit a pretty big issue, *the DC component seems to push for the very low values of $\lambda_{1,2}$, even though visually, they're not necessarily the best*. For this, Brian suggested to look at the zero-point of k -space. This may be a good idea regardless, as it seems like there may be a light roll off where there shouldn't be in the raw image data – it can be seen in the real part of the data.

As of 06/29/15, this has been fixed, however the original data has been put in `/projects/muisjes/asalerno/CS/data/raw_test_data/data`. The issue that existed was that for a 256×256 dataset, the $(k_x, k_y) = (0, 0)$ was not in the proper location.

It seems like the data is now well normalized and there actually isn't a drop in the mean of the data after applying the CS routine. The *max* seems to drop, but the mean actually increases by a fraction of a percent. The code `/home/asalerno/Documents/CompressedSensing/demo_TVxfm.m` has been edited in order to include the ability to normalize either by max or mean.

06/30/15

Something interesting that I noted was that the first iteration performs the most “clean up” of the data. If one were to run the code, looking at the data after each loop, I'd expect that they'd see a major jump after the first run, then the differences to taper off exponentially after that – like a power function.

7.3 Noise Analysis

By taking a small subset of the image where the noise exists, I ran a simple noise analysis on the data. The noise analysis that was done is akin to the one done in MBP 1024 for the MRI lab. In fact, it uses a slightly modified version of that code.

When this analysis was done, it appears as if the noise is Rayleigh distributed *after* the Compressed Sensing algorithm, regardless of the parameter choice (under $\lambda_{1,2} \in [10^{-3}, 1]$ – the same analysis is to be done for $\lambda_{1,2} \in [1, 10^4]$). To test this, I ran the analysis on each image coming out of

the parameter optimization algorithm. The noise distributions were fit to a Lorentzian, Gaussian, and Rayleigh function, with the highest value of an adjusted R^2 value taken to be the “best fit”.

Doing the same analysis for the data before the compressed sensing algorithm, we find that, again, we exist in a low SNR regime of MRI where the data fits nicely to a Rayleigh function. Thus, I believe it is OK to utilize our normal statistics on the data.

7.4 Overly Smoothed Results with $\lambda_1 = \lambda_2 = 0$ – 06/30/15

Figure 11: Overly smoothed result. Comparisons of when there is no compressed sensing, to when $\lambda_1 = \lambda_2 = 0$, and again to $\lambda_1, \lambda_2 \neq 0$

Some weird results are occurring when $\lambda_1 = \lambda_2 = 0$. Essentially, we’d expect that we end up with something that is very close to the original dataset, as can be seen in the figure above. However, this isn’t the case – we are obtaining an overly smoothed result that looks almost nothing like the original `im_dc` variable as we would expect. The question here is why.

Looked into whether it was the XFM that was causing it. Upon some quick inspection, I can safely say it wasn’t. On the first iteration, when it goes through the line $x = x + t * dx$, the major difference comes. Somewhere in this line is a major change that occurs. The question, however, is exactly what is happening...

Apparently, FTXFMtx looks quite a lot like a brain (as it should as it’s apparently just the k -space image of the brain)... This could be where we’re getting the utter craziness. Is this what

is building our first guess? Perhaps our first guess is a very smoothed out version of the data that we have.

As of 06/30/15, we've figured out that the issue actually arises from what exactly is being taken as `im_res` and `data`. Because of this, we have many different angles that we have to look at.

7.5 Analysis of the code – What's our starting point? – 07/02/15 -

07/02/15

Looked at the code and found how similar the data looks right from the get go – the basis that we are using is `data`, even though the data that we are feeding into it is `im_dc`, which is density compensated. Thus, after the first iteration, it's being pushed towards the non-density-compensated case. Thus, even when $\lambda_1 = \lambda_2 = 0$, we get this nicely smoothed image.

07/03/15

Looking into this a little more, I want to see exactly when the data tanks towards the `data` variable, i.e. when there is no density compensation.